

Step 1: Column Pruning

What is Column Pruning?

Column Pruning is the optimization where the query engine reads only the specific columns requested in a `SELECT` statement, ignoring all others.

Why Columnar Formats Enable It

In a columnar format like Parquet, each column is stored in its own separate block or set of files. The storage layer has a "map" (metadata) that tells the engine exactly where the `amount` column starts and ends on the disk.

Why Row-Based Formats Cannot

Row-based formats (CSV, JSON, Avro) store data as a sequence of records. To reach the 6th column in a row, the computer must physically scan through the first 5 columns to find the delimiter. Even if you only need 2 columns, you pay the **I/O cost** for the entire row.

Example Scenario

Table: `orders` (9 columns: `order_id`, `customer_id`, `order_date`, `region`, `product_id`, `amount`, `discount`, `tax`, `total_amount`) **Query:** `SELECT order_id, amount FROM orders;`

- **Total Columns:** 9
 - **Columns Needed:** 2
 - **Columnar (Parquet):** 2 columns read.
 - **Row-based (CSV):** 9 columns read (the whole file).
 - **I/O Savings:** ~78% reduction in data transferred from disk to memory.
-

Step 2: Predicate Pushdown

What is Predicate Pushdown?

Predicate Pushdown is the optimization where the filtering logic (the `WHERE` clause) is "pushed" down to the storage layer. The engine evaluates whether a block of data is relevant *before* actually opening and reading it.

How Parquet Metadata Helps

Parquet files are divided into **Row Groups**. For every row group, Parquet stores metadata (statistics) in the file footer, including:

- **Min/Max values** for every column.
- **Null counts**.

Why it Reduces Data Scanned

If you filter `WHERE order_date = '2026-02-21'`, the engine checks the Min/Max metadata for each Row Group. If a group's `order_date` range is `2026-01-01` to `2026-02-15`, the engine **skips that entire block** without reading a single row.

Analysis

- **Predicate:** `order_date = '2024-03-01'`
 - **Metadata Help:** Engine looks at the `order_date` min/max for each block.
 - **Data Skipped:** Every row group where '2024-03-01' is outside the min/max range.
 - **I/O Reduction:** Massive. If you have 3 years of data, you might skip 99% of the file.
-

Step 3: Query Scenario Analysis

Query: `SELECT region, SUM(amount) FROM sales WHERE order_date = '2026-02-21' GROUP BY region;`

Column Analysis

- **Table Columns:** `order_id, customer_id, order_date, region, product_id, amount, discount`.
- **Needed Columns:** `order_date` (for filter), `region` (for grouping), `amount` (for sum).
- **Columns Read:** 3 (`order_date, region, amount`).
- **Columns Skipped:** 4 (`order_id, customer_id, product_id, discount`).

Optimization Breakdown

Step	Action	Benefit
1. Pushdown	Check footer metadata for <code>order_date</code> min/max.	Skips non-matching row groups entirely.
2. Pruning	Locate only <code>order_date</code> , <code>region</code> , and <code>amount</code> offsets.	Skips 4 unused columns within matching groups.
3. Reading	Pull only the 3 columns from matching groups into RAM.	Minimizes network/disk traffic.
4. Processing	Aggregate <code>amount</code> by <code>region</code> in CPU.	Operates on a tiny, pre-filtered subset of data.

Step 4: Optimization Summary

The Synergy

Parquet provides the **structure** (Columnar), Column Pruning provides the **focus** (Selective Read), and Predicate Pushdown provides the **filter** (Skipping). Together, they allow a query engine to act like a laser rather than a floodlight.

Scale Impact Example

- **Table Size:** 1 Billion Rows, 20 Columns (Assume 100GB total size).
- **Query Needs:** 2 columns (10% of width), filters to 1% of rows (Date filter).
- **CSV Reading:** 100 GB (Must read everything).
- **Parquet Reading:** ~100 MB (10% of the columns \times 1% of the rows).
- **Result: 1000x improvement** in I/O and significant cost savings in cloud environments like AWS (S3/Athena).