

Step 1-4: Order Dataset Analysis

Dataset: `orders(order_id, customer_id, order_date, amount)`

Candidate Shard Key Evaluation

Shard Key	Pros	Cons	Use Case Fit
<code>order_id</code>	High cardinality; perfectly even distribution if using UUIDs.	Sequential IDs create a hot shard on the most recent server. Queries by <code>customer_id</code> must hit all shards (scatter-gather).	Poor for this workload (customer-centric).
<code>customer_id</code>	Co-locates all orders for one customer on one shard. Fast single-shard reads.	Risk of data skew if one "Power User" has millions of orders.	Best Fit for OLTP/Customer portals.
<code>order_date</code>	Excellent for time-range analytics and archiving old data.	Massive hot shard problem: 100% of today's writes go to the "Today" shard.	Best for logs/analytics only.

Step 5: Other Dataset Evaluations

1. Application Logs

- **Best Shard Key:** `timestamp` (Range-based).
- **Reasoning:** Logs are rarely updated; we mostly care about "what happened in the last hour." Range sharding allows us to drop or archive an entire "Month" shard instantly.

2. User Profiles

- **Best Shard Key:** `user_id` (Hash-based).
- **Reasoning:** Even distribution is vital here. We want to ensure that as we grow to 100 million users, the load is spread perfectly across all nodes.

3. IoT Sensor Data

- **Best Shard Key:** `device_id + timestamp` (Composite).
 - **Reasoning:** We usually want to see a specific device's history. Sharding by `device_id` ensures one device's data is together, but adding `timestamp` allows us to sub-partition by time for better performance.
-

Step 6: Reflection Questions

1. **Why is the choice irreversible?** To change a shard key, you must create a new cluster, define the new key, and move **all** data from the old cluster to the new one while the system is live. This causes massive "rebalancing" traffic and potential downtime.
 2. **Analytics vs. OLTP?** OLTP wants to find **one user's** data instantly (Single Shard). Analytics wants to sum **everyone's** data over a time period (Multiple Shards/Range).
 3. **Cross-shard queries?** These require a "Coordinator" node to send the query to every shard, wait for all results, and merge them. It is the "slowest link" problem—if one shard is slow, the whole query is slow.
-

Step 7: E-commerce Product Catalog Design

Product Catalog Sharding Design

```
## Chosen Shard Key  
**category_id** (Hash-based)
```

Reasoning

Most shoppers browse by category (e.g., "Electronics"). By sharding on `category_id`, all products in that category live on the same shard, making "Browse" queries extremely fast.

Query Routing

- **By Category:** Routed to a single shard.
- **By Product ID:** Requires a Global Index or a scatter-gather query across shards.
- **Search:** Handled by a separate search engine (like Elasticsearch), not the sharded DB.

Edge Cases

- **Skew:** If "Clothing" has 80% of products, we can use a composite key `category_id + subcategory_id` to break the hot shard into smaller pieces.

Step 8: Comparison Matrix

Strategy	Distribution	Query Routing	Complexity	Best For
Hash-based	Excellent	Single shard for Key	Low	OLTP, User Profiles
Range-based	Can be uneven	Best for ranges	Medium	Time-series, Logs
Composite	Balanced	Flexible	High	IoT, Complex SaaS
Directory	Perfectly Flexible	Dynamic	Very High	Multi-tenant apps